

# CSE 584 – MACHINE LEARNING: TOOLS AND ALGORITHMS

## HOMEWORK-2

Divya Navuluri  
PSU ID: 903661629

---

### **ABSTRACT THAT PROVIDES A HIGH-LEVEL OVERVIEW OF THE REINFORCEMENT LEARNING CODE'S PURPOSE AND THE OVERALL PROCESS:**

I have taken the below code from the online GitHub. The Reinforcement learning code covers the implementation of a Deep Q-Network with TensorFlow and Keras that solves reinforcement learning tasks by letting an agent learn how to select optimal policies while interacting with an environment. The DQN is built upon off-policy Q-learning, where it uses a deep neural network in place of tabular Q-value approximation for the selection of actions. It contains two major components in the network: a primary model for training and making action decisions, and a target model to give stable Q-value estimates during training.

The DQN class is initiated with several parameters, state dimensions, action space, discount factor often known as gamma, exploration rate known as epsilon, and learning rate. There are two hidden layers in the model architecture, both with ReLU as their activation functions. This enables the network to learn complex patterns in the input state space. Action selection is based on an epsilon-greedy policy, which forms a good trade-off between exploration and exploitation for optimal learning.

The training process consists of replaying experiences from a memory buffer, where mini-batches are sampled to update the weights of the main model, based on predicted Q-values and corresponding rewards. Then, it comes to a soft update mechanism for synchronizing the target model with the primary one in order to ensure stability during training. The exploration rate grows more conservatory over time in order to gradually shift the agent's behavior from exploration toward exploitation.

Besides, it includes functions to save model architecture and weights that would allow the saving and reusability of the trained DQN model. Overall, this implementation serves as a foundational framework for applying deep reinforcement learning techniques to various decision-making problems.

## CORE SECTION OF THE REINFORCEMENT LEARNING IMPLEMENTATION

### CODE:

```
1 import numpy as np
2 from keras.models import Sequential
3 from keras.models import model_from_json
4 from keras.layers import Dense, Activation
5 from keras import optimizers
6 from keras import backend as K
7 import tensorflow as tf
8 from random import random, randrange
9
10 # 1.DQN Class Initialization and Model Creation
11 # Deep Q Network (DQN) class implementing off-policy Q-learning with deep neural networks
12 class DQN:
13
14     def __init__(self,
15                 input_dim,          # Number of inputs for the DQN network
16                 action_space,       # Size of action space
17                 gamma=0.99,         # Discount factor for future rewards, balancing immediate and future rewards
18                 epsilon=1,          # Initial epsilon
19                 epsilon_min=0.01,   # Minimum epsilon value to maintain some exploration in the long term
20                 epsilon_decay=0.999, # Decay rate of epsilon after each action taken to reduce exploration over time
21                 learning_rate=0.00025, # Learning rate for training the neural network
22                 tau=0.125,          # Soft update factor for the target network, controls how fast it tracks the main model
23                 model=None,         # Placeholder for the main DQN model
24                 target_model=None,  # Placeholder for the target DQN model used for stable Q-learning targets
25                 sess=None):        # TensorFlow session
26
27         # Initialize parameters for the DQN
28         self.input_dim = input_dim # Set the input dimension for the DQN
29         self.action_space = action_space # Set the size of the action space
30         self.gamma = gamma # Set the discount factor
31         self.epsilon = epsilon # Initialize epsilon for exploration
32         self.epsilon_min = epsilon_min # Set the minimum epsilon for exploration
33         self.epsilon_decay = epsilon_decay # Set the decay rate for epsilon
34         self.learning_rate = learning_rate # Set the learning rate for model training
35         self.tau = tau # Set tau for soft updates of the target network
36
37         # Create the main and target DQN models
38         self.model = self.create_model() # Main DQN model for training and action selection
39         self.target_model = self.create_model() # Target DQN model to stabilize learning
40
41         # TensorFlow configuration for GPU memory optimization
42         config = tf.compat.v1.ConfigProto() # Create a TensorFlow configuration object
43         config.gpu_options.allow_growth = True # Allow the GPU memory to grow as needed to prevent allocation issues
44         self.sess = tf.compat.v1.Session(config=config) # Create a TensorFlow session with the specified configuration
45         K.set_session(self.sess) # Set the created session as the backend session for Keras
46         self.sess.run(tf.compat.v1.global_variables_initializer()) # Initialize global variables in TensorFlow
47
48     def create_model(self):
49         model = Sequential() # Initialize a sequential model for easy stacking of layers
50         model.add(Dense(300, input_dim=self.input_dim)) # First hidden layer with 300 neurons, taking the input dimension
51         model.add(Activation('relu')) # ReLU activation function introduces non-linearity
52
53         model.add(Dense(300)) # Second hidden layer with 300 neurons
54         model.add(Activation('relu')) # ReLU activation function for the second hidden layer
55
56         model.add(Dense(self.action_space)) # Output layer with a number of neurons equal to the action space size
57         model.add(Activation('linear')) # Linear activation function for producing Q-values
58
59         # Define the optimizer and compile the model
60         sgd = optimizers.SGD(lr=self.learning_rate, decay=1e-6, momentum=0.95) # Stochastic Gradient Descent optimizer
61         model.compile(optimizer=sgd, loss='mse') # Compile the model with Mean Squared Error loss function for Q-learning
62
63         return model # Return the compiled model for use in training and action selection
64
```

```

64
65
66 # 2. Action Selection
67 def act(self, state):
68     # Selects an action based on an epsilon-greedy policy
69     a_max = np.argmax(self.model.predict(state.reshape(1, len(state)))) # Get action with maximum Q-value
70     if random() < self.epsilon: # epsilon-greedy decision with probability epsilon, choose random action
71         a_chosen = randrange(self.action_space) # Randomly select an action
72     else:
73         a_chosen = a_max # Choose action with the highest Q-value otherwise
74     return a_chosen # Return the selected action
75
76 # 3. replay Method and training the model
77 def replay(self, samples, batch_size):
78     # Trains the DQN model on a mini-batch of experiences
79     inputs = np.zeros((batch_size, self.input_dim)) # Initialize batch input array
80     targets = np.zeros((batch_size, self.action_space)) # Initialize batch target Q-values array
81
82     for i in range(batch_size): # Loop through each sample in the batch
83         state = samples[0][i, :] # Get current state from samples
84         action = samples[1][i] # Get action taken from samples
85         reward = samples[2][i] # Get reward received for action
86         new_state = samples[3][i, :] # Get resulting state from samples
87         done = samples[4][i] # Check if this state is terminal
88
89         inputs[i, :] = state # Set input as current state for batch
90         targets[i, :] = self.target_model.predict(state.reshape(1, len(state))) # Predict current Q-values for state
91
92         if done: # If terminal state, set target as immediate reward
93             targets[i, action] = reward # Assign the reward to the Q-value of the action taken
94         else:
95             q_future = np.max(self.target_model.predict(new_state.reshape(1, len(new_state)))) # Predict future max Q-value
96             targets[i, action] = reward + Q_future * self.gamma # Set target Q-value using discounted future reward
97
98     # Train the model on the current batch of inputs and target Q-values
99     loss = self.model.train_on_batch(inputs, targets)
100
101 # 4. target_train Method
102 def target_train(self):
103     # Updates the target model with weights from the main model using a soft update
104     weights = self.model.get_weights() # Get weights from main DQN model
105     target_weights = self.target_model.get_weights() # Get weights from target DQN model
106     for i in range(len(target_weights)): # Update each layer's weights
107         target_weights[i] = weights[i] * self.tau + target_weights[i] * (1 - self.tau) # Soft update rule
108     self.target_model.set_weights(target_weights) # Set updated weights in target model
109
110 # 5. Epsilon Update
111 def update_epsilon(self):
112     # Reduces the epsilon to encourage exploitation over time
113     self.epsilon = self.epsilon * self.epsilon_decay # Decay epsilon
114     self.epsilon = max(self.epsilon_min, self.epsilon) # Ensure epsilon does not fall below minimum value
115
116 # 6. Model Saving
117 def save_model(self, path, model_name):
118     # Saves the model structure and weights to disk for later usage
119     model_json = self.model.to_json() # Serialize model to JSON format
120     with open(path + model_name + ".json", "w") as json_file: # Open a file to save the model structure
121         json_file.write(model_json) # Save JSON structure
122     self.model.save_weights(path + model_name + ".h5") # Save model weights in HDF5 format
123     print("Saved model to disk") # Print confirmation of save

```

**1. DQN Class Initialization and Model Creation:** The DQN class implements a Deep Q-Network, which uses off-policy Q-learning with neural networks. It initializes parameters like input dimensions, action space, discount factor, and exploration settings, and creates both a primary model for choosing actions and a target model for stable training. This is a two-hidden-layer sequence model, compiled with Stochastic Gradient Descent and mean squared error loss to have effective learning.

**2. Action Selection:** Act method adopts an epsilon-greedy strategy in picking an action, which balances exploration and exploitation. It predicts the Q-values for all the actions, picks a random action if the generated number is less than epsilon, or otherwise picks the action with the highest Q-value. This way, while making effective decisions, the agent can still explore new actions.

**3. replay method and training the model:** The replay method trains the model using mini-batches from experience replay memory. It prepares the input states and target Q-values, retrieves the necessary information for each sample, and calculates target values based on the immediate rewards and future maximum Q-values. The model is then updated with efficient mini-batch training, returning the loss for this iteration.

**4. target\_train Method:** The target\_train method performs a soft update of the target model weights. It pulls the weights from the main and target models then sets the target model weights to be a mix of its current weights and those of the main model, controlled by a specified tau. parameter.

**5. Epsilon Update:** The update\_epsilon method decreases the rate of exploration epsilon in such a way that, increasingly, the more the agent is trained, the exploitation of learned Q-values is favored. It decays epsilon by a factor, while never letting it drop below a minimum value, trying to encourage more confident action selections over time.

**6. Model Saving:** The save\_model method allows for saving the architecture along with the weights of DQN. This serializes the model architecture in JSON format and saves it, along with the model weights in HDF5 format, confirming that saving has been successful for later retrieval and use.